

# Smoothing and Regression Splines

*Ignasi Mañé, Antoni Company & Chiara Barbi*

*19 de marzo de 2019*

## Introduction

The file `bikes.Washington.Rdata` contains information on the bike-sharing rental service in Washington D.C., USA, corresponding to years 2011 and 2012. This file contains only one data frame, `bikes`, with 731 rows (one for each day of years 2011 and 2012, that was a leap year) and 9 columns:

- `instant`: row index, going from 1 to 731.
- `yr`: year (0: 2011, 1:2012).
- `dayyr`: day of the year (from 1 to 365 for 2011, and from 1 to 366 for 2012).
- `weekday`: day of the week (0 for Sunday, 1 for Monday, ..., 6 for Saturday).
- `workingday`: if day is neither weekend nor holiday is 1, otherwise is 0.
- `temp`: temperature in Celsius.
- `hum`: humidity in %.
- `windspeed`: wind speed in miles per hour.
- `cnt`: count of total rental bikes. In this exam we consider this variable as continuous.

In the following chunk we will call the libraries used throughout the assignment.

```
library(ggplot2)
library(splines)
library(tidyverse)
library(splines)
```

## Nonparametric regression of `cnt` as a function of `instant`

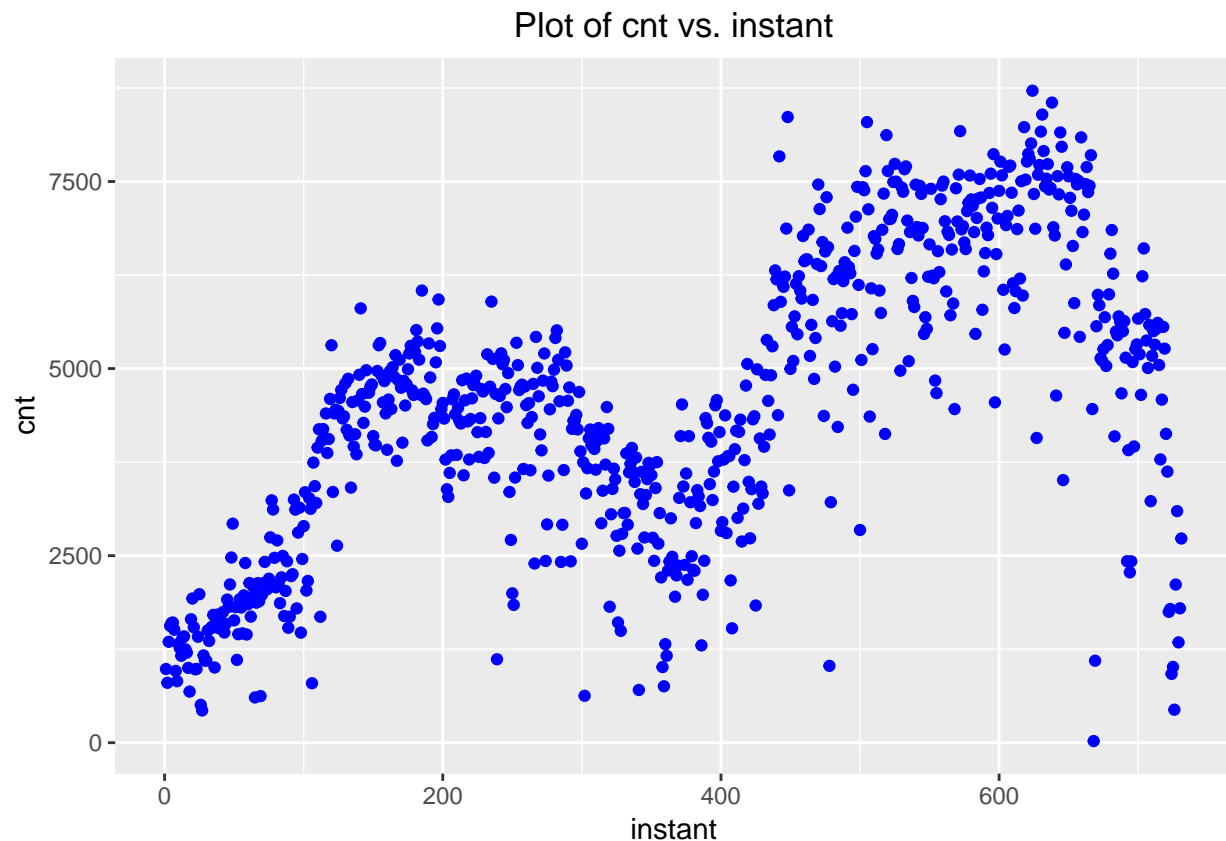
1. Consider the nonparametric regression of `cnt` as a function of `instant`. Estimate the regression function  $m(\text{instant})$  of `cnt` as a function of `instant` using a cubic regression splines estimated with the R function `smooth.splines` and choosing the smoothing parameter by Generalized Cross Validation.

First, we will proceed by uploading the data to the R environment and defining two auxiliary parameters for the variables `cnt` and `instant`.

```
load("bikes.Washington.RData")
cnt <- bikes$cnt
instant <- bikes$instant
```

Once we have uploaded the data we will proceed with a simple representation of the cnt as a function of instant to see the behavior of the parameters.

```
ggplot(bikes)+geom_point(aes(instant,cnt), col='blue')+
  ggtitle(label="Plot of cnt vs. instant")+
  theme(plot.title = element_text(hjust = 0.5))
```



Once we have seen the data, we will proceed with the questions for this exercise.

**a) Which is the value of the chosen penalty parameter  $\lambda$ ?**

We have chosen  $\lambda$  by applying the generalized cross-validation for the data to analyse, to obtain its value we applied the function `smooth.spline()` and we imposed `cv=FALSE` to do the generalized cross-validation.

```
s.gcv <- smooth.spline(x = instant,y = cnt,cv=FALSE)
```

Where the optimal value for  $\lambda$  is:

```
s.gcv$lambda
```

```
## [1] 1.005038e-07
```

b) Which is the corresponding equivalent number of degrees of freedom  $df$ ?

The function `smooth.spline()` used before also gives us the value for the degrees of freedom, where  $df$  will be:

```
s.gcv$df
```

```
## [1] 93.34091
```

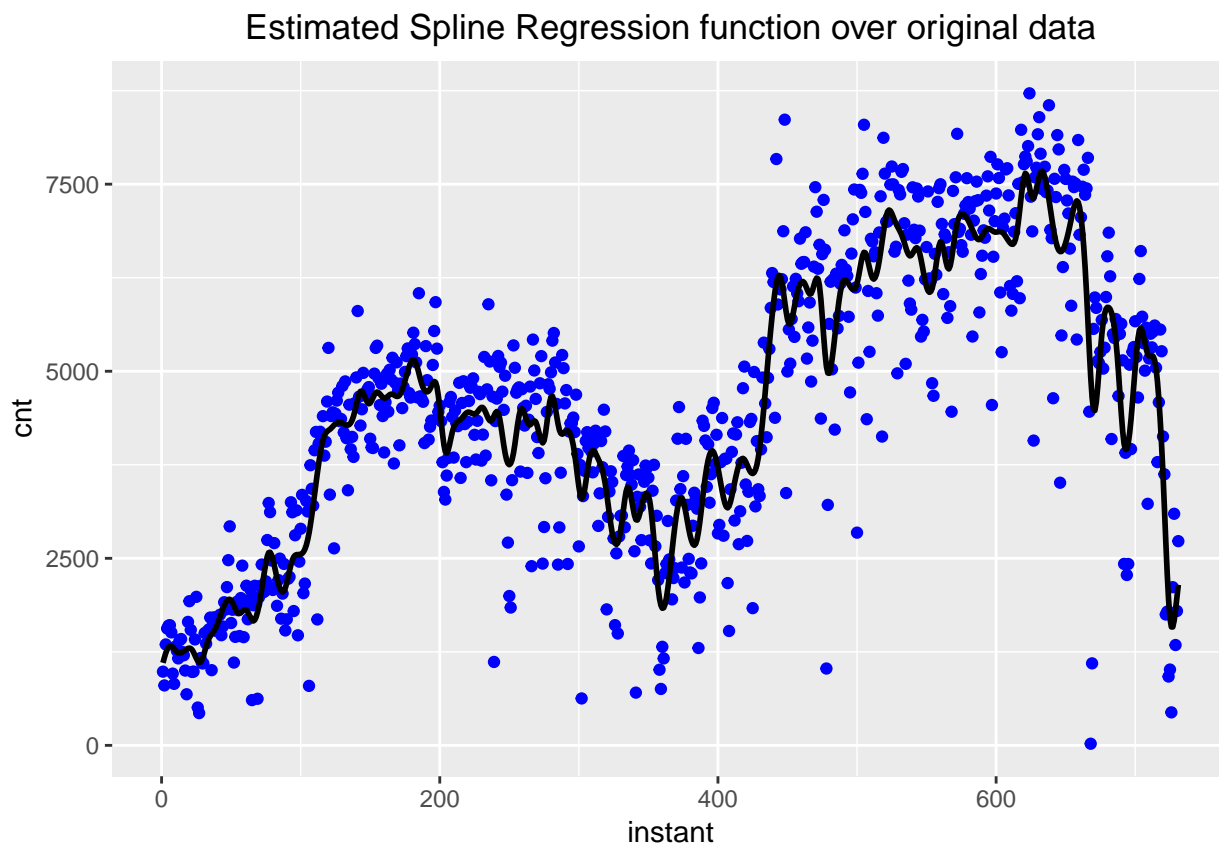
c) How many knots have been used?

```
s.gcv$fit$nk - 2
```

```
## [1] 134
```

d) Give a graphic with the scatter plot and the estimated regression function  $m(\text{instant})$ .

```
ggplot(bikes)+geom_point(aes(instant,cnt), col='blue')+  
  ggtitle(label="Estimated Spline Regression function over original data")+  
  theme(plot.title = element_text(hjust = 0.5))+  
  geom_line(data = as.data.frame(s.gcv$x,s.gcv$y),  
            aes(x =s.gcv$x,y =s.gcv$y),color="black",size=1)
```



e) Estimate now  $m(\text{instant})$  by unpenalized regression splines combining the R functions `bs` and `lm`, using the knots `my.knots <- quantile(instant, ((1:n.knots)-.5)/n.knots)` where `n.knots` is the previous value of `df` minus 4.

```
n.knots <- s.gcv$df-4
my.knots <- quantile(instant, ((1:n.knots)-0.5)/n.knots)
my.knots <- my.knots[-c(1,length(my.knots))]
```

Once the data is prepared, we can proceed computing the cubic B-Spline basis used for the posterior `lm()` of `cnt` as a function of the basis.

```
basis <- bs(x=instant,knots = my.knots,intercept = TRUE,degree = 3)
```

Then, we finally proceed to compute the `lm()`.

```
lm.bs <- lm(cnt~basis-1)
```

If we summarize the results from the `lm()` we can see that:

```
summary(lm.bs)
```

```
##
## Call:
## lm(formula = cnt ~ basis - 1)
```

```
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4856.9  -360.8   112.5   502.8  2487.2
##
## Coefficients:
##      Estimate Std. Error t value Pr(>|t|)
## basis1      965.9      686.5   1.407 0.159899
## basis2     1732.7      968.8   1.789 0.074166 .
## basis3      923.5      935.2   0.988 0.323760
## basis4     1453.0      700.5   2.074 0.038448 *
## basis5      995.0      692.0   1.438 0.150982
## basis6     1544.7      689.7   2.240 0.025443 *
## basis7     2087.8      689.0   3.030 0.002542 **
## basis8     1673.5      688.8   2.430 0.015389 *
## basis9     1932.5      688.8   2.806 0.005171 **
## basis10    1365.0      688.7   1.982 0.047925 *
## basis11    3449.6      688.7   5.009 7.10e-07 ***
## basis12    1159.2      688.7   1.683 0.092857 .
## basis13    3193.8      688.7   4.637 4.28e-06 ***
## basis14    1993.5      688.7   2.894 0.003928 **
## basis15    3863.9      688.7   5.610 3.01e-08 ***
## basis16    4217.9      688.7   6.124 1.59e-09 ***
## basis17    4443.4      688.7   6.452 2.19e-10 ***
## basis18    4029.1      688.7   5.850 7.84e-09 ***
## basis19    5052.5      688.7   7.336 6.69e-13 ***
## basis20    4322.1      688.7   6.275 6.43e-10 ***
## basis21    4996.5      688.7   7.255 1.17e-12 ***
## basis22    4278.5      688.7   6.212 9.41e-10 ***
## basis23    5335.0      688.7   7.746 3.73e-14 ***
## basis24    4773.7      688.7   6.931 1.02e-11 ***
## basis25    5053.7      688.7   7.338 6.61e-13 ***
## basis26    4176.0      688.7   6.063 2.28e-09 ***
## basis27    3992.6      688.7   5.797 1.06e-08 ***
## basis28    4526.3      688.7   6.572 1.03e-10 ***
## basis29    4549.3      688.7   6.605 8.35e-11 ***
## basis30    4147.4      688.7   6.022 2.91e-09 ***
## basis31    4999.5      688.7   7.259 1.13e-12 ***
## basis32    2851.3      688.7   4.140 3.94e-05 ***
## basis33    5162.2      688.7   7.495 2.21e-13 ***
## basis34    3974.4      688.7   5.771 1.23e-08 ***
## basis35    4184.0      688.7   6.075 2.13e-09 ***
## basis36    4630.0      688.7   6.723 3.96e-11 ***
## basis37    4272.6      688.7   6.204 9.90e-10 ***
## basis38    3457.3      688.7   5.020 6.71e-07 ***
```

## basis39	3425.8	688.7	4.974	8.44e-07	***
## basis40	4657.2	688.7	6.762	3.07e-11	***
## basis41	1820.0	688.7	2.642	0.008431	**
## basis42	3899.1	688.7	5.661	2.27e-08	***
## basis43	2490.0	688.7	3.615	0.000324	***
## basis44	4104.1	688.7	5.959	4.19e-09	***
## basis45	1926.3	688.7	2.797	0.005316	**
## basis46	1503.9	688.7	2.184	0.029356	*
## basis47	4373.9	688.7	6.351	4.06e-10	***
## basis48	1965.0	688.7	2.853	0.004469	**
## basis49	3334.5	688.7	4.842	1.62e-06	***
## basis50	4566.3	688.7	6.630	7.14e-11	***
## basis51	2511.7	688.7	3.647	0.000287	***
## basis52	3968.0	688.7	5.761	1.30e-08	***
## basis53	3850.8	688.7	5.591	3.34e-08	***
## basis54	3206.4	688.7	4.656	3.93e-06	***
## basis55	6318.6	688.7	9.174	< 2e-16	***
## basis56	6038.3	688.7	8.767	< 2e-16	***
## basis57	5775.4	688.7	8.385	3.22e-16	***
## basis58	5836.6	688.7	8.474	< 2e-16	***
## basis59	7165.0	688.7	10.403	< 2e-16	***
## basis60	3724.1	688.7	5.407	9.05e-08	***
## basis61	6961.8	688.7	10.108	< 2e-16	***
## basis62	5387.9	688.7	7.823	2.14e-14	***
## basis63	7172.5	688.7	10.414	< 2e-16	***
## basis64	5624.7	688.7	8.167	1.69e-15	***
## basis65	7268.4	688.7	10.553	< 2e-16	***
## basis66	7196.9	688.7	10.450	< 2e-16	***
## basis67	6207.6	688.7	9.013	< 2e-16	***
## basis68	7133.1	688.7	10.357	< 2e-16	***
## basis69	5331.6	688.7	7.741	3.86e-14	***
## basis70	7333.9	688.7	10.648	< 2e-16	***
## basis71	5895.0	688.7	8.559	< 2e-16	***
## basis72	7839.7	688.7	11.383	< 2e-16	***
## basis73	6370.6	688.7	9.250	< 2e-16	***
## basis74	6943.0	688.7	10.081	< 2e-16	***
## basis75	7146.2	688.7	10.376	< 2e-16	***
## basis76	6139.9	688.7	8.915	< 2e-16	***
## basis77	7796.8	688.7	11.321	< 2e-16	***
## basis78	7059.9	688.7	10.251	< 2e-16	***
## basis79	8078.5	688.7	11.729	< 2e-16	***
## basis80	6662.0	688.7	9.673	< 2e-16	***
## basis81	6013.4	688.7	8.731	< 2e-16	***
## basis82	8782.5	688.7	12.752	< 2e-16	***
## basis83	4281.4	688.7	6.217	9.16e-10	***

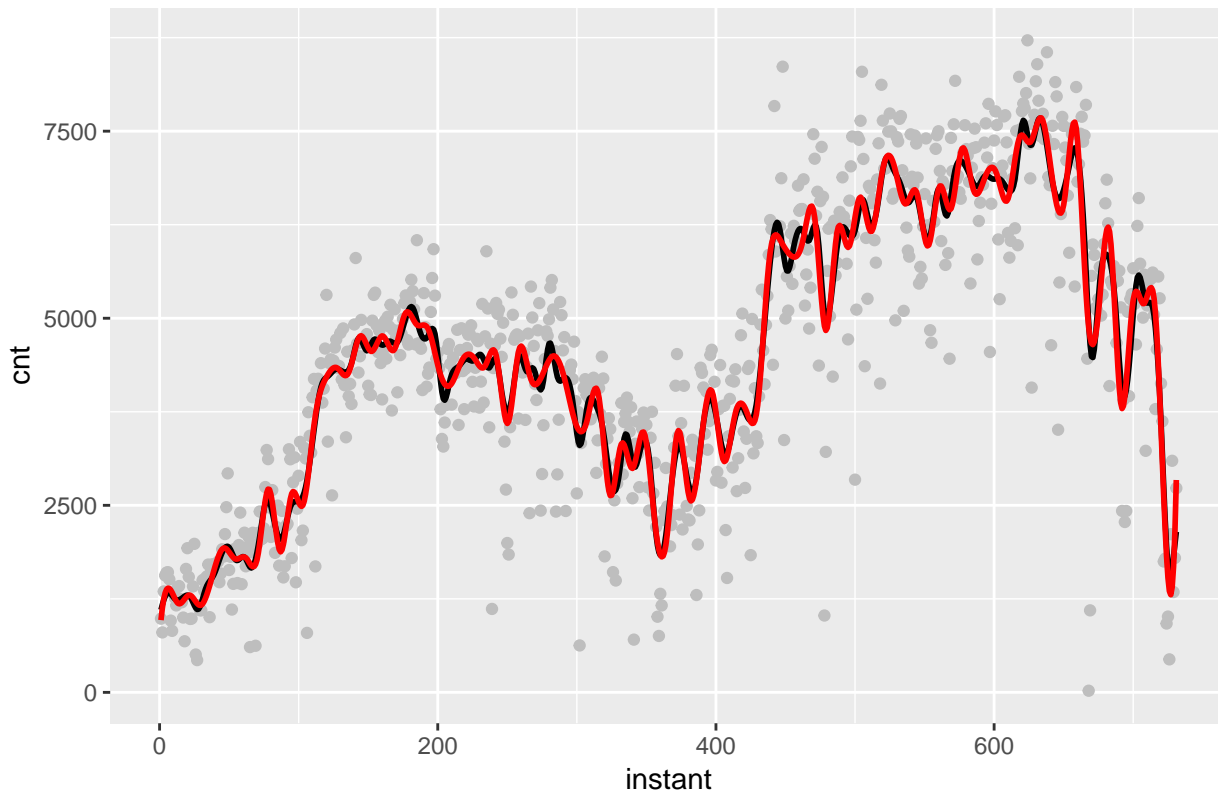
```
## basis84    4713.6      688.6    6.846 1.79e-11 ***
## basis85    7412.7      688.1   10.772 < 2e-16 ***
## basis86    2374.4      686.7    3.458 0.000581 ***
## basis87    6015.9      681.7    8.825 < 2e-16 ***
## basis88    4729.4      665.1    7.111 3.09e-12 ***
## basis89    6648.8      918.9    7.236 1.33e-12 ***
## basis90   -1205.7      957.2   -1.260 0.208288
## basis91    2838.5      650.9    4.361 1.51e-05 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 883.6 on 640 degrees of freedom
## Multiple R-squared:  0.9716, Adjusted R-squared:  0.9675
## F-statistic: 240.3 on 91 and 640 DF,  p-value: < 2.2e-16
```

The model adjusted is highly significant, and only a few bases could be discarded for having a p-value under the 0.05.

f ) Give a graphic with the scatter plot and the two estimated regression functions.

```
ggplot(bikes)+geom_point(aes(instant,cnt), col='grey')+
  ggtitle(label="Comparision of Penalized vs Unpenalized Regression Splines")+
  theme(plot.title = element_text(hjust = 0.5))+
  geom_line(data = as.data.frame(s.gcv$x,s.gcv$y),aes(x =s.gcv$x,y =s.gcv$y),color="black")+
  geom_line(aes(x=instant,y=lm.bs$fitted.values),color="red",size=1)
```

### Comparison of Penalized vs Unpenalized Regression Splines



## Nonparametric logistic regression using splines

2. Nonparametric logistic regression using splines with a IRWLS procedure. The script `IRWLS_logistic_regression.R` includes the definition of the function `logistic.IRWLS.splines` performing non-parametric logistic regression using splines with a IRWLS procedure. The basic syntax is the following: `logistic.IRWLS.splines(x=..., y=..., x.new=..., df=..., plts=TRUE)` where the arguments are the explanatory variable `x`, the 0-1 response variable `y`, the vector `x.new` of new values of variable `x` where we want to predict the probability of `y` being 1 given that `x` is equal to `x.new`, the equivalent number of parameters (or model degrees of freedom) `df`, and the logical `plts` indicating if plots are desired or not. Define a new variable `cnt.5000` taking the value 1 for days such that the number of total rental bikes is larger than or equal to 5000, on 0 otherwise.

a) Use the function `logistic.IRWLS.splines` to fit the non-parametric binary regression `cnt.5000` as a function of the temperature, using `df=6`. In which range of temperatures is  $\Pr(\text{cnt} \geq 5000 | \text{temp})$  larger than 0,5?

First we have to upload the script “`IRWLS_logistic_regression.R`” in order to get the function `logistic.IRWLS.splines()`



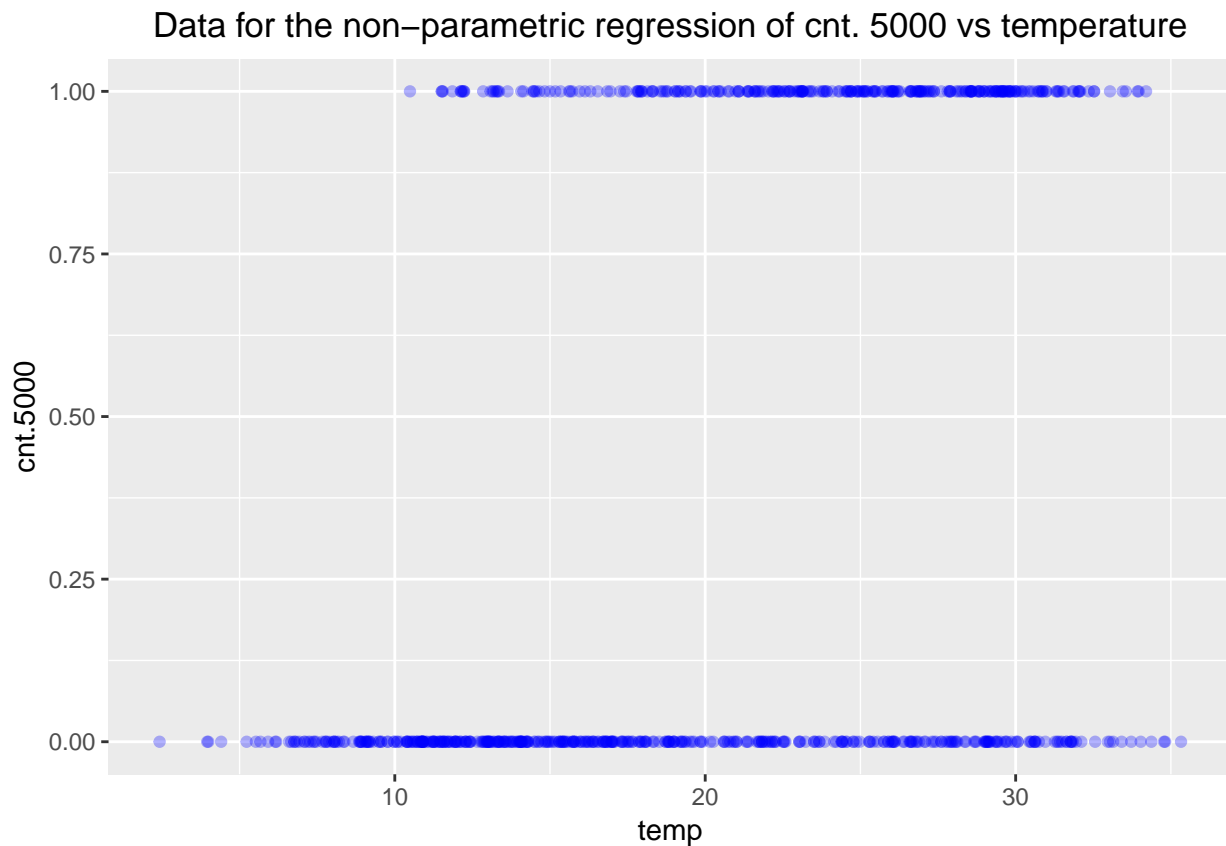
```
source("IRWLS_logistic_regression.R")
```

Once we have the function uploaded, we also have to define the parameters for the non-parametric regression, where:

```
temp <- bikes$temp #x axis  
#and the response variable is defined as follows:  
cnt.5000 = ifelse(bikes$cnt>=5000,1,0)
```

We will represent a plot in order to see the data for the non-parametric regression:

```
ggplot(as.data.frame(temp,cnt.5000))+geom_point(aes(temp,cnt.5000), col='blue',alpha=0.3)  
ggtitle(label="Data for the non-parametric regression of cnt. 5000 vs temperature")+th
```



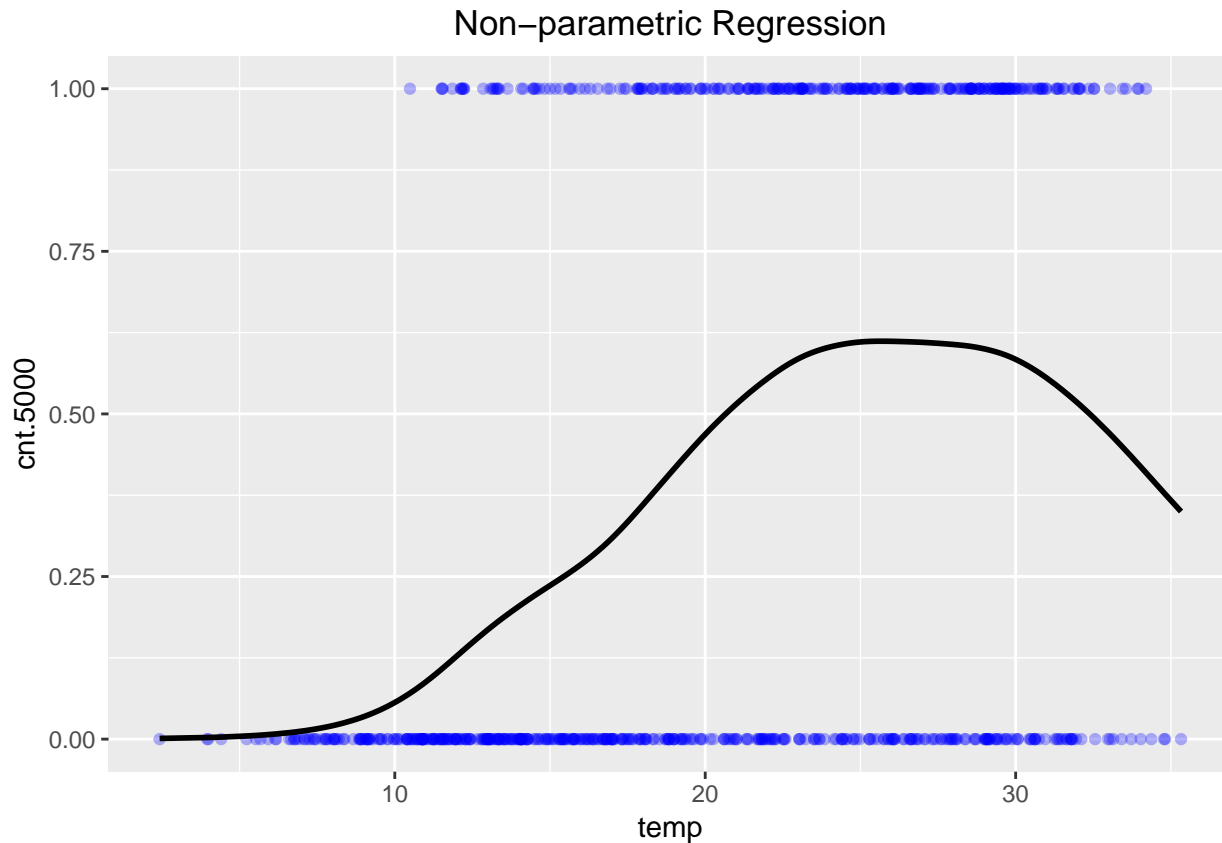
Once we have seen the behavior of the data, we can proceed with the non-parametric regression with the logistic function with  $df=6$ .

```
log.reg <- logistic.IRWLS.splines(y=cnt.5000,x = temp, x.new = temp, df = 6, plts = F)
```

Where if we analyse the results of the regression we will see that:

```
ggplot(as.data.frame(temp,cnt.5000))+  
geom_point(aes(temp,cnt.5000), col='blue',alpha=0.3)+
```

```
ggtitle(label="Non-parametric Regression")+
theme(plot.title = element_text(hjust = 0.5))+
geom_line(data = as.data.frame(temp,log.reg$fitted.values),
          aes(temp,log.reg$predicted.values),color='black',size=1)
```



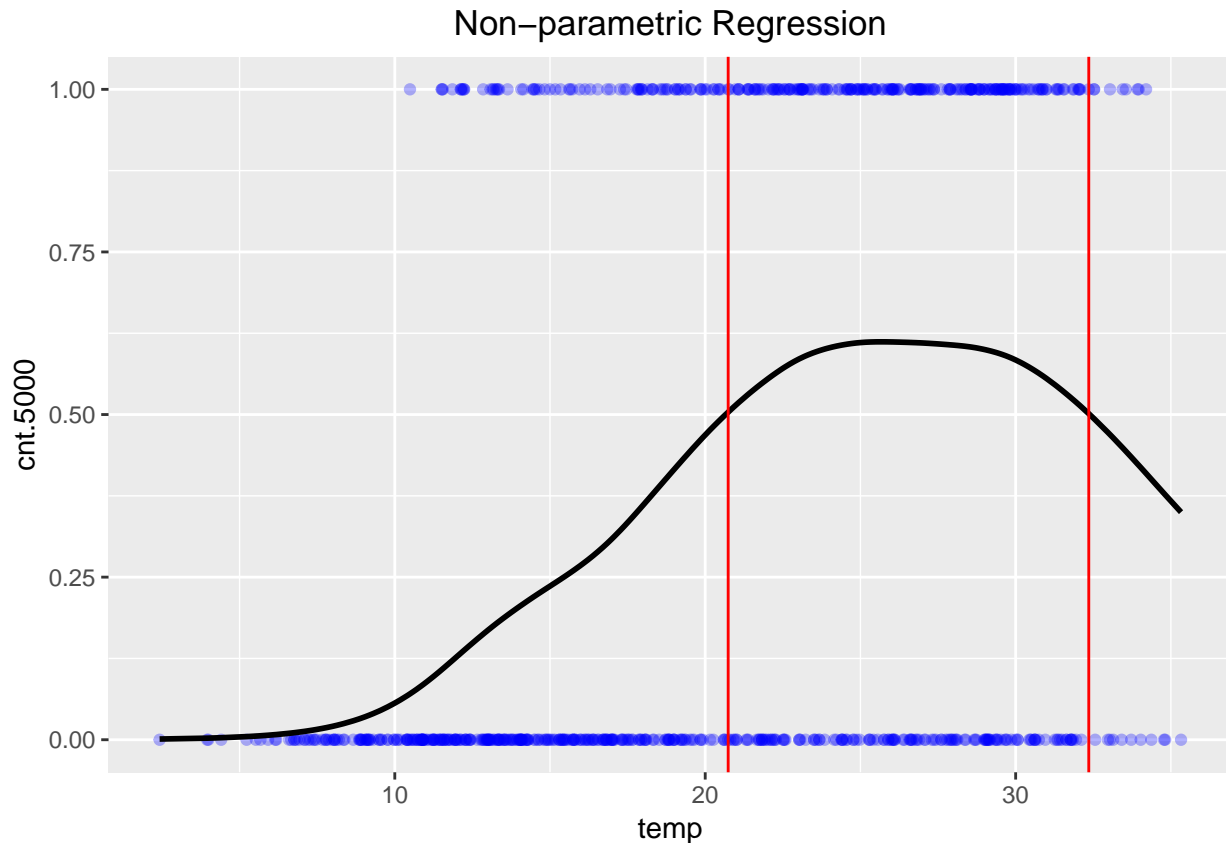
Where the range of the percentage of temperature which are larger than  $\text{cnt.5000} \geq 0.5$  is the following:

```
range.0.5 <- range(temp[log.reg$fitted.values>=0.5])
range.0.5
```

```
## [1] 20.73915 32.35585
```

If we analyse the results over the previous plot:

```
ggplot(as.data.frame(temp,cnt.5000))+geom_point(aes(temp,cnt.5000), col='blue',alpha=0.5)+
ggtitle(label="Non-parametric Regression")+theme(plot.title = element_text(hjust = 0.5))+
geom_line(data = as.data.frame(temp,log.reg$fitted.values),
          aes(temp,log.reg$predicted.values),color='black',size=1)+
geom_vline(xintercept = range.0.5[1],color='red')+
geom_vline(xintercept = range.0.5[2],color='red')
```



b) (Optional) Choose the parameter  $df$  by  $k$ -fold cross validation with  $k = 5$  and using  $df.v = 3:15$  as the set of possible values for  $df$ .

For this exercise we will use the same parameters as before, but, we also have to define a sequence of the degrees of freedom in order to obtain the optimal  $df^*$  by using the 5-fold cross-validation method.

```
df.s <- seq(3,15,by=1)
k <- 5
```

Once we have defined the degrees of freedom and the method of  $k$ -fold cross-validation, we will proceed to compute cross-validation in order to obtain the optimal value for the degrees of freedom.

To compute the  $k$ -fold cross-validation we have implemented the following code, which measures the accuracy in the predictions (the same as for the assignment of the delivery ROC) of the regression as a reference.

```
set.seed(1994)
accuracy <- NULL

for (i in 1:length(df.s)){

  #folds definition
```

```

folds <- as.numeric(sample(rep(1:k,length.out = length(temp))))
acc.df.iteration <- NULL

for (j in 1:k ){

  #validation data
  x.val <- temp[which(folds == j)]
  y.val <- cnt.5000[which(folds == j)]
  #training data
  x.train <- temp[which(folds != j)]
  y.train <- cnt.5000[which(folds != j)]

  #regression estimation for each df value
  log.reg.cv <- logistic.IRWS.splines(x = x.train,y = y.train,x.new = x.val,df = df.s[j])

  #accuracy in the regression
  acc.table <- table(y.val,ifelse(log.reg.cv$predicted.values>0.5,1,0))
  acc.df.iteration[j] <- (acc.table[1,1]+acc.table[2,2])/sum(acc.table)
}

accuracy[i] <- mean(acc.df.iteration)
}

```

We plot the results of the iteration process:

```

ggplot(as.data.frame(df.s,accuracy))+
  geom_point(aes(df.s,accuracy))+
  geom_line(data=as.data.frame(df.s,accuracy),aes(df.s,accuracy))+
  ggtitle(label="Accuracy in the prediction")+
  theme(plot.title = element_text(hjust = 0.5))+
  geom_vline(xintercept = df.s[which.max(accuracy)],color='red')

```

